
▮ New Dataset: Retail Transactions

▮ Sample Data (Create this as `retail_data.csv`)

```
TransactionID, Customer, City, Product, Category, Quantity, UnitPrice, TotalPrice, TransactionDate, PaymentMode
T1001, Ali, Mumbai, Laptop, Electronics, 1, 70000, 70000, 2024-01-15, Card
T1002, Neha, Bangalore, Tablet, Electronics, 2, 30000, 60000, 2024-01-20, UPI
T1003, Ravi, Hyderabad, Desk, Furniture, 1, 15000, 15000, 2024-02-10, Net Banking
T1004, Zoya, Delhi, Chair, Furniture, 4, 5000, 20000, 2024-02-12, Card
T1005, Karan, Mumbai, Phone, Electronics, 1, 50000, 50000, 2024-02-15, Card
T1006, Farah, Delhi, Mouse, Electronics, 3, 1000, 3000, 2024-02-18, Cash
```

▮ Task Set – PySpark Hands-On (No DLT)

▮ Basics

1. Load `retail_data.csv` into a PySpark DataFrame and display schema.
 2. Infer schema as False – then manually cast columns.
-

▮ Data Exploration & Filtering

3. Filter transactions where `TotalPrice > 40000`.
 4. Get unique cities from the dataset.
 5. Find all transactions from "Delhi" using `.filter()` and `.where()`.
-

▮ Data Manipulation

6. Add a column `DiscountedPrice = TotalPrice - 10%`.
 7. Rename `TransactionDate` to `TxnDate`.
 8. Drop the column `UnitPrice`.
-

▮ Aggregations

9. Get total sales by city.
 10. Get average unit price by category.
 11. Count of transactions grouped by PaymentMode.
-

▮ Window Functions

12. Use a window partitioned by City to rank transactions by `TotalPrice`.
 13. Use lag function to get previous transaction amount per city.
-

▮ Joins

14. Create a second DataFrame `city_region`:

```
City, Region
Mumbai, West
Delhi, North
```

```
Bangalore, South  
Hyderabad, South
```

15. Join with main DataFrame and group total sales by Region.

▮ Nulls and Data Cleaning

- 16. Introduce some nulls and replace them with default values.
 - 17. Drop rows where Quantity is null.
 - 18. Fill null PaymentMode with "Unknown".
-

▮ Custom Functions

- 19. Write a UDF to label orders:

```
def label_order(amount):  
    if amount > 50000: return "High"  
    elif amount >= 30000: return "Medium"  
    else: return "Low"
```

Apply this to classify TotalPrice .

▮ Date & Time

- 20. Extract year, month, and day from TxnDate .
 - 21. Filter transactions that happened in February.
-

▮ Union & Duplicate Handling

- 22. Duplicate the DataFrame using union() and remove duplicates.
-